


● Pre-seed 2026

Minds & Models

The **who-layer** of retail measurement.

Retail measurement has always known what sold.
We add who bought, and join the two.



A model of how people shop physical retail, that brands can ask questions of.

02 · WHERE THE MONEY IS SPENT BLIND

Every source a brand buys today reads the same thing.
The receipt. Never the shopper.

Sell-out and POS data

Every basket, every store, every week. The category's backbone for forty years.

- No person attached to the basket

Consumer panels

A few thousand recruited households, weighted and projected to a national market.

- Self-reported, and months late

Loyalty card data

One retailer, its own card holders, its own definition of a segment.

- Blind outside that chain, and to every non-buyer

THE SIZE OF THE BLIND SPOT

Nine in ten FMCG euros are spent in a store where nobody knows who is buying.

E-commerce solved this twenty years ago. The shelf never did.

Ecommerce takes 7.2% of global grocery spend (Kantar Worldpanel). Everything else happens in a store.

Every store already produces two data sets. **Nobody has ever joined them.**

WHO

The shopper

- 100+ visual attributes per shopper
- Read on the edge, nothing identifiable leaves the store

Observed, not declared. Not scraped, not generated

WHAT

The receipt

- Every line, every basket, every store
- The retailer's own POS feed, already clean

Exists already. The half the incumbents sell

↓ MATCHED AT THE VISIT ↓

THE RESULT

Who bought what.

A model of how people shop physical retail, queried in plain language. **The join is the product, and nobody else has it.**

04 · WHO PAYS FOR WHAT

The side that has the data does not pay. The side that needs it does.

SUPPLY

The retailer

€0 Pays nothing, ever

— GIVES

- **The cameras it already owns**, plus the POS feed for the same stores. No hardware from us, no capex from them.
- **A perpetual aggregate license**, anonymized and chain-masked, signed once per chain. This is the asset we accumulate.

— GETS

- **A revenue share on data it never had to collect.** No budget line means no CFO and no procurement committee.

DEMAND

The brands

€40k → €170k Per category, per year

— GIVES

- **A measurement budget that already exists**, locked every October and today spent on half the picture.
- **No exclusivity.** The same category read is sold to every brand in that category. Syndication is the multiplier.

— GETS

- **Who bought, matched to what sold.** The answer their current stack cannot produce at any price.

WHAT A BRAND PAYS, PER CATEGORY, PER YEAR

40k EUR

Founding seat, first chain live



80k EUR

The read is continuous in one chain



170k EUR

The read spans two or more chains



Value is non-linear. The old way scaled with installations. Here each chain deepens every read already running, so **revenue rises without a new sale.**

You are spending on a shopper you have never seen. **Ask who they are.**

Ask anything, or @ a segment to name who you mean



Ask @Dawn Trades

Ask #energy drinks

Ask #convenience stores

Ask @Young Streetwear

Ask #deodorant

Segments are discovered in the data, not authored by a marketer.

PROMOTION EFFICIENCY

“Who pays full price for deodorant, and who waits for the promotion?”

→ Stop funding discount on the segment that would have paid anyway

CHANNEL STRATEGY

“How does our shopper mix differ between convenience and drugstores?”

→ Different pack size, price tier and listing argument per channel

AUDIENCE BUILDING

“Which competitor's shopper looks most like ours, and where do they overlap?”

→ Aim the switching campaign at the audience that is actually winnable

TRY IT chat.mindsmodels.ai →

Tested with market leaders.

TELCO RETAIL



O2

PET SPECIALIST



Super zoo

COSMETICS

MANUFAKTURA

Manufaktura

SPORTS NUTRITION



GymBeam

HEALTH FOOD



Vilgain

FUEL AND CONVENIENCE



EuroOil

RETAIL REAL ESTATE



CPI Property Group

VISUAL DATA POINTS

18,830,280

Interpreted shopper signal, not raw footage

POS DATA POINTS MATCHED

484,120

Each one joined to the visit that produced it


Every store-month adds to the same thing: **an owned, interpreted corpus of real-world visual shopper data**, tied to verified purchases. It is the one part of this business that cannot be rebuilt by spending more.

07 · THE PLAN AND WHAT IT IS WORTH

Software multiples are compressing.


Owned data is repricing the other way.

ENDING ARR, EUR MILLIONS



18.1M EUR of ARR on 7.1M EUR of capital, profitable in 2031 at a 22.7% December run-rate.

WHERE ASSETS LIKE THIS CHANGE HANDS, ON REVENUE



WHAT THAT MULTIPLE IS ACTUALLY PRICING

Not the software. An owned corpus of real-world visual shopper data, held under a perpetual cross-chain license signed one retailer at a time. That signature is the one thing more capital cannot accelerate.

TOP-DOWN CHECK · IS THE MARKET ACTUALLY BIG ENOUGH

\$18bn

Global consumer measurement, as the incumbent sizes its own market

NIQ Form S-1, SEC, 2025

→

23,000

Companies already buying it from that one incumbent, at ~\$173k a year each

NIQ Form S-1, average implied

→

137

Brand seats is our entire 2031 plan. 0.6% of one incumbent's client list.

Minds & Models financial model

Meet the team.



Ondřej Šantora

Co-Founder

4+ years in GTM scaling of AI products



ondrej@mindsmodels.ai · +420 720 571 736



Josef Němeček

Co-Founder

Research in algorithmic personalization



josef@mindsmodels.ai · +420 737 138 638

THE TEAM'S COLLECTIVE EXPERIENCE



Consumer Intelligence @ NielsenIQ

Recommender systems @ Alza

Product & GTM @ SpaceKnow

Computer vision

AI infrastructure

500+ academic pages

09 · WHAT WE ARE RAISING AND WHAT IT PROVES

A 500k EUR pre-seed buys one thing: proof that this works on real deployments.

01

Retailers sign.

Chains move from pilot to a **perpetual data mandate**. This is the single thing every euro of the plan rests on, and it has never been tested at signature.

02

Brands pay.

Category reads we have already delivered convert into **annual seats, billed in advance**. A study proves interest. A renewal proves a market.

03

It holds at scale.

Two categories running **continuously across live chains**, not one window in one store. This is where a demo becomes a product.

04

The ladder is real.

The first seat **repricing itself** as a second chain goes live. If that happens once, the whole non-linear model stops being a slide and becomes an observation.

10 months

of runway, enough to put all four to the test

On evidence

the next round is priced on signed revenue, not on a plan

21 Nov 2026

rolling close, targeted

APPENDIX · UNIT ECONOMICS

Why the margin recovers and the seat gets more valuable.

METRIC	2027	2029	2031	READ
Gross margin	71.3%	64.4%	74.5%	Dips while new markets are opened, then recovers as slots fill with seats
Net revenue retention	260%	168%	110%	Computed from slots crossing the price gates, not assumed. Early years sit on a base of one to five slots and are comparatively meaningless
ARPU per brand seat	70.3k	105.2k	132.6k	Against a \$50k floor and a \$75k to \$500k enterprise band at the incumbent
LTV / CAC	6.0×	7.0×	8.4×	Improves as founder-led selling is replaced by a repeatable motion into an installed base
CAC payback	9.9 mo	8.6 mo	7.2 mo	Fully loaded: all sales, retailer BD and marketing payroll plus programs
Live slots · seats each	5 · 1.4	22 · 2.6	48 · 2.9	A slot is one category in one channel. A handful of brands per category are large enough to buy
Headcount	17	48	74	245k EUR of ARR per head at the end of the plan

Retailer revenue share falls from 29% to 22% of revenue across the plan, because the cross-chain aggregated read carries no share: no single chain can claim it.

APPENDIX · THE PLAN AND THE CAPITAL

Six years, and what each round pays for.

EUR	2026	2027	2028	2029	2030	2031
Ending ARR	94k	494k	1.99M	5.58M	11.51M	18.12M
Revenue recognized	18k	270k	1.23M	3.60M	8.34M	15.20M
EBITDA	(87k)	(842k)	(2.73M)	(2.63M)	(1.48M)	2.31M
Closing cash	491k	2.24M	3.86M	2.07M	1.97M	5.82M
Headcount	8	17	34	48	64	74

ROUND	WHEN	AMOUNT	POST	AT EXIT
Angel, done	Q2 2026	80k	3.5M	1.4%
Pre-seed	Nov 2026	500k	6.5M	5.3%
Seed	Sep 2027	2.5M	12M	18.1%
Series A	Nov 2028	4.0M	30M	13.3%

SEED

Opens Poland, pulls DACH forward, funds the second and third chain, and produces the first cross-retailer aggregated read. ARR is about 302k EUR at that point, so the round is priced on evidence.

SERIES A

Buys the United States, at a one-off entry cost of 600k EUR. The largest measurement market in the world, and where the buyers of this company are.

7.12M EUR raised across all sources. Founders hold **55.7%** at exit, the option pool 6.2%.

APPENDIX · THE EXIT LADDER

What each multiple assumes the buyer thinks they are buying.

MULTIPLE	EXIT VALUE	ON 2031 REVENUE	THIS ROUND RETURNS	WHAT THAT MULTIPLE PRICES
2.3×	41.7M	2.7×	4.4×	Undifferentiated store analytics. A footfall incumbent was acquired at about this level, which is the market's verdict on analytics sold without owned rights. This is the floor if the rights work fails.
4×	72.5M	4.8×	7.7×	A rented asset. A license the retailer can terminate is a service, and services in this category price here.
8×	145.0M	9.5×	15.3×	An owned rights portfolio, at the low end of the measurement band.
10×	181.2M	11.9×	19.1×	The plan. Midpoint of the measurement band, priced against the buyer that already sells the what-sold half and cannot assemble the who half itself.
12×	217.4M	14.3×	23.0×	The same asset in a competitive process, or with more markets live than this plan carries inside its horizon.

The spread from 2.3× to 12× is roughly 176M EUR of exit value, and it turns on whether the license is perpetual and cross-chain rather than on the technology. That is why the data-processing terms sit in the assumptions of the financial model rather than in a legal annex.

APPENDIX · WHAT BREAKS IT

The four questions worth asking, answered before you ask them.

What if the Series A never happens?

The United States never opens and 2031 ARR lands near half the base case. The company still never runs out of cash. **The Series A does not buy survival, it buys the American half of the exit.**

What if churn is three times the assumption?

Tested at 15% and 25% against a modeled 8%. The 25% case is what sized the Series A at 4.0M EUR. At that size every tested case stays funded.

Is profitability too far out?

EBITDA turns positive in 2031 at 18.1M EUR of ARR on 7.1M EUR of total capital. The full-year margin of 15.2% averages loss-making months with profitable ones. **December run-rate is 22.7%.**

What is genuinely unproven?

Rights breadth. Everything above assumes we sign perpetual aggregate mandates chain by chain, and that is the constraint this round exists to remove. The ladder prices the case where we do not, and it still returns 4.4×